

Exercises

- 1.1 Explain the difference between fast computing and real-time computing.
- 1.2 What are the main limitations of the current real-time kernels for the development of critical control applications?
- 1.3 Discuss the features that a real-time system should have for exhibiting a predictable timing behavior.
- 1.4 Describe the approaches that can be used in a real-time system to handle peripheral I/O devices in a predictable fashion.
- 1.5 Which programming restrictions should be used in a programming language to permit the analysis of real-time applications? Suggest some extensions that could be included in a language for real-time systems.

Exercises

- 2.1 Give the formal definition of a schedule, explaining the difference between preemptive and non-preemptive scheduling.
- 2.2 Explain the difference between periodic and aperiodic tasks, and describe the main timing parameters that can be defined for a real-time activity.
- 2.3 Describe a real-time application as a number of tasks with precedence relations, and draw the corresponding precedence graph.
- 2.4 Discuss the difference between static and dynamic, on-line and off-line, optimal, and heuristic scheduling algorithms.
- 2.5 Provide an example of domino effect, caused by the arrival of a task J^* , in a feasible set of three tasks.

Exercises

- 3.1 Check whether the Earliest Due Date (EDD) algorithm produces a feasible schedule for the following task set (all tasks are synchronous and start at time $t = 0$):

	J_1	J_2	J_3	J_4
C_i	4	5	2	3
D_i	9	16	5	10

- 3.2 Write an algorithm for finding the maximum lateness of a task set scheduled by the EDD algorithm.
- 3.3 Draw the full scheduling tree for the following set of non-preemptive tasks and mark the branches that are pruned by the Bratley's algorithm.

	J_1	J_2	J_3	J_4
a_i	0	4	2	6
C_i	6	2	4	2
D_i	18	8	9	10

- 3.4 On the scheduling tree developed in the previous exercise find the path produced by the Spring algorithm using the following heuristic function: $H = a + C + D$. Then find a heuristic function that produces a feasible schedule.
- 3.5 Given seven tasks, A , B , C , D , E , F , and G , construct the precedence graph from the following precedence relations:

$$\begin{array}{ll}
 A \rightarrow C & \\
 B \rightarrow C & B \rightarrow D \\
 C \rightarrow E & C \rightarrow F \\
 D \rightarrow F & D \rightarrow G
 \end{array}$$

Then, assuming that all tasks arrive at time $t = 0$, have deadline $D = 25$, and computation times 2, 3, 3, 5, 1, 2, 5, respectively, modify their arrival times and deadlines to schedule them by EDF.

Exercises

- 4.1 Verify the schedulability and construct the schedule according to the RM algorithm for the following set of periodic tasks:

	C_i	T_i
τ_1	2	6
τ_2	2	8
τ_3	2	12

- 4.2 Verify the schedulability and construct the schedule according to the RM algorithm for the following set of periodic tasks:

	C_i	T_i
τ_1	3	5
τ_2	1	8
τ_3	1	10

- 4.3 Verify the schedulability and construct the schedule according to the RM algorithm for the following set of periodic tasks:

	C_i	T_i
τ_1	1	4
τ_2	2	6
τ_3	3	10

- 4.4 Verify the schedulability under RM of the following task set:

	C_i	T_i
τ_1	1	4
τ_2	2	6
τ_3	3	8

- 4.5 Verify the schedulability under EDF of the task set shown in Exercise 4.4, and then construct the corresponding schedule.
- 4.6 Verify the schedulability under EDF and construct the schedule of the following task set:

	C_i	D_i	T_i
τ_1	2	5	6
τ_2	2	4	8
τ_3	4	8	12

- 4.7 Verify the schedulability of the task set described in Exercise 4.6 using the Deadline-Monotonic algorithm. Then construct the schedule.

Exercises

- 5.1 Compute the maximum processor utilization that can be assigned to a Sporadic Server to guarantee the following periodic tasks under RM:

	C_i	T_i
τ_1	1	5
τ_2	2	8

- 5.2 Compute the maximum processor utilization that can be assigned to a Deferrable Server to guarantee the task set illustrated in Exercise 5.1.
- 5.3 Together with the periodic tasks illustrated in Exercise 5.1, schedule the following aperiodic tasks with a Polling Server having maximum utilization and intermediate priority.

	a_i	C_i
J_1	2	3
J_2	7	1
J_3	17	1

- 5.4 Solve the same scheduling problem described in Exercise 5.3, with a Sporadic Server having maximum utilization and intermediate priority.
- 5.5 Solve the same scheduling problem described in Exercise 5.3, with a Deferrable Server having maximum utilization and highest priority.
- 5.6 Using a Sporadic Server with capacity $C_s = 2$ and period $T_s = 5$, schedule the following tasks:

periodic tasks

	C_i	T_i
τ_1	1	4
τ_2	2	6

aperiodic tasks

	a_i	C_i
J_1	2	2
J_2	5	1
J_3	10	2

Exercises

- 6.1 Compute the maximum processor utilization that can be assigned to a Dynamic Sporadic Server to guarantee the following periodic tasks, under EDF:

	C_i	T_i
τ_1	2	6
τ_2	3	9

- 6.2 Together with the periodic tasks illustrated in Exercise 6.1, schedule the following aperiodic tasks with a Dynamic Sporadic Server with $C_s = 2$ and $T_s = 6$.

	a_i	C_i
J_1	1	3
J_2	5	1
J_3	15	1

- 6.3 Solve the same scheduling problem described in Exercise 6.2 with a Total Bandwidth Server having utilization $U_s = 1/3$.
- 6.4 Solve the same scheduling problem described in Exercise 6.2 with a Constant Bandwidth Server with $C_s = 2$ and $T_s = 6$.
- 6.5 Solve the same scheduling problem described in Exercise 6.2 with an Improved Total Bandwidth Server with $U_s = 1/3$, which performs only one shortening step.
- 6.6 Solve the same scheduling problem described in Exercise 6.2 with the optimal Total Bandwidth Server (TB*).
- 6.7 Consider the following set of periodic tasks:

	C_i	T_i
τ_1	4	10
τ_2	4	12

After defining two Total Bandwidth Servers, TB_1 and TB_2 , with utilization factors $U_{s1} = 1/10$ and $U_{s2} = 1/6$, construct the EDF schedule in the case in which two aperiodic requests $J_1(a_1 = 1, C_1 = 1)$ and $J_2(a_2 = 9, C_2 = 1)$ are served by TB_1 , and two aperiodic requests $J_3(a_3 = 2, C_3 = 1)$ and $J_4(a_4 = 6, C_4 = 2)$ are served by TB_2 .

Exercises

- 7.1 Verify whether the following task set is schedulable by the Rate-Monotonic algorithm. Apply the processor utilization approach first, and then the Response Time Analysis:

	C_i	T_i	B_i
τ_1	4	10	5
τ_2	3	15	3
τ_3	4	20	0

- 7.2 Consider three periodic tasks τ_1 , τ_2 , and τ_3 (having decreasing priority), which share three resources, A , B , and C , accessed using the Priority Inheritance Protocol. Compute the maximum blocking time B_i for each task, knowing that the longest duration $D_i(R)$ for a task τ_i on resource R is given in the following table (there are no nested critical sections):

	A	B	C
τ_1	2	0	2
τ_2	2	3	0
τ_3	3	2	5

- 7.3 Solve the same problem described in Exercise 7.2 when the resources are accessed by the Priority Ceiling Protocol.
- 7.4 For the task set described in Exercise 7.2, illustrate the situation produced by RM + PIP in which task τ_2 experiences its maximum blocking time.
- 7.5 Consider four periodic tasks τ_1 , τ_2 , τ_3 , and τ_4 (having decreasing priority), which share five resources, A , B , C , D , and E , accessed using the Priority Inheritance Protocol. Compute the maximum blocking time B_i for each task, knowing that the longest duration $D_i(R)$ for a task τ_i on resource R is given in the following table (there are no nested critical sections):

	A	B	C	D	E
τ_1	2	5	9	0	6
τ_2	0	0	7	0	0
τ_3	0	3	0	7	13
τ_4	6	0	8	0	10

- 7.6 Solve the same problem described in Exercise 7.5 when the resources are accessed by the Priority Ceiling Protocol.
- 7.7 For the task set described in Exercise 7.5, illustrate the situation produced by RM + PIP in which task τ_2 experiences its maximum blocking time.
- 7.8 Consider three tasks τ_1 , τ_2 , and τ_3 , which share three multi-unit resources, A , B , and C , accessed using the Stack Resource Policy. Resources A and B have three units, whereas C has two units. Compute the ceiling table for all the resources based on the following task characteristics:

	D_i	μ_A	μ_B	μ_C
τ_1	5	1	0	1
τ_2	10	2	1	2
τ_3	20	3	1	1

1

SOLUTIONS TO THE EXERCISES

SOLUTIONS FOR CHAPTER 1

- 1.1 Fast computing tends to minimize the average response time of computation activities, whereas real-time computing is required to guarantee the timing constraints of each task.
- 1.2 The main limitations of the current real-time kernels are mainly due to the fact that they are developed to minimize runtime overhead (hence functionality) rather than offering support for a predictable execution. For example, short interrupt latency is good for servicing I/O devices, but introduces unpredictable delays in task execution for the high priority given to the interrupt handlers. Scheduling is mostly based on fixed priority, thus explicit timing constraints cannot be specified on tasks. No specific support is usually provided for periodic tasks, and no aperiodic service mechanism is available for handling event-driven activities. Access to shared resources is often realized through classical semaphores, which are efficient, but prone to priority inversion, if no protocol is implemented for entering critical sections. Finally, no temporal protection or resource reservation mechanism is usually available in current real-time kernels for coping with transient overload conditions, so a task executing too much may introduce unbounded delays on the other tasks.
- 1.3 A real-time kernel should allow the user to specify explicit timing constraints on application tasks and support a predictable execution of real-time activities with specific real-time mechanisms, including scheduling, resource management, synchronization, communication, and interrupt handling. In critical real-time

systems, predictability is more important than high performance, and often an increased functionality can only be reached at the expense of a higher runtime overhead. Other important features that a real-time system should have include maintainability, fault-tolerance, and overload management.

- 1.4 Three approaches can be used. The first one is to disable all external interrupts, letting application tasks to access peripheral devices through polling. This solution gives great programming flexibility and reduces unbounded delays caused by the driver execution, but it is characterized by a low processor efficiency on I/O operations, due to the busy wait.

A second solution is to disable interrupts and handle I/O devices by polling through a dedicated periodic kernel routine, whose load can be taken into account through a specific utilization factor. As in the previous solution, the major problem of this approach is due to the busy wait, but the advantage is that all hardware details can be encapsulated into a kernel routine and do not need to be known by the application tasks. An additional overhead is due to the extra communication required between application tasks and kernel routines for exchanging I/O data.

A third approach is to enable interrupts but limit the execution of interrupt handlers as much as possible. In this solution, the interrupt handler activates a device handler, which is a dedicated task that is scheduled (and guaranteed) by the kernel as any other application task. This solution is efficient and minimizes the interference caused by interrupts.

- 1.5 The restrictions that should be used in a programming language to permit the analysis of real-time applications should limit the variability of execution times. Hence, a programmer should avoid using dynamic data structures, recursion, and all high level constructs that make execution time unpredictable. Possible language extensions should be aimed at facilitating the estimation of worst-case execution times. For example, a language could allow the programmer to specify the maximum number of iterations in each loop construct, and the probability of taking a branch in conditional statements.

SOLUTIONS FOR CHAPTER 2

- 2.1 A schedule is formally defined as a step function $\sigma : \mathbf{R}^+ \rightarrow \mathbf{N}$ such that $\forall t \in \mathbf{R}^+, \exists t_1, t_2$ such that $t \in [t_1, t_2)$ and $\forall t' \in [t_1, t_2) \sigma(t) = \sigma(t')$. For any $k > 0$, $\sigma(t) = k$, means that task J_k is executing at time t , while $\sigma(t) = 0$

means that the CPU is idle. A schedule is said to be *preemptive* if the running task can be arbitrarily suspended at any time to assign the CPU to another task according to a predefined scheduling policy. In a preemptive schedule, tasks may be executed in disjointed interval of times. In a non-preemptive schedule, a running task cannot be interrupted and therefore it proceeds until completion.

- 2.2 A periodic task consists of an infinite sequence of identical jobs, that are regularly activated at a constant rate. If ϕ_i is the activation time of the first job of task τ_i , the activation time of the k th job is given by $\phi_i + (k - 1)T_i$, where T_i is the task period. Aperiodic tasks also consist of an infinite sequence of identical jobs; however, their activations are not regular. An aperiodic task where consecutive jobs are separated by a minimum interarrival time is called a sporadic task. The most important timing parameters defined for a real-time task are
- the *arrival time* (or *release time*), that is, the time at which a task becomes ready for execution;
 - the *computation time*, that is, the time needed by the processor for executing the task without interruption;
 - the *absolute deadline*, that is, the time before which a task should be completed to avoid damage to the system;
 - the *finishing time*, that is, the time at which a task finishes its execution;
 - the *response time*, that is, the difference between the finishing time and the release time: $R_i = f_i - r_i$;
- 2.3 A real-time application consisting of tasks with precedence relations is shown in Section 2.2.2.
- 2.4 A static scheduler is one in which scheduling decisions are based on fixed parameters, assigned to tasks before their activation. In a dynamic scheduler, scheduling decisions are based on dynamic parameters that may change during system evolution. A scheduler is said to be *off line* if it is pre-computed (before task activation) and stored in a table. In an on-line scheduler, scheduling decisions are taken at runtime when a new task enters the system or when a running task terminates. An algorithm is said to be *optimal* if it minimizes some given cost function defined over the task set. A common optimality criterion for real-time system is related to feasibility. Then, a scheduler is optimal whenever it can find a feasible schedule, if there exists one. Heuristic schedulers use a heuristic function to search for a feasible schedule, hence it is not guaranteed that a feasible solution is found.
- 2.5 An example of domino effect is shown in Figure ??.

SOLUTIONS FOR CHAPTER 3

- 3.1 To check whether the EDD algorithm produces a feasible schedule, tasks must be ordered with increasing deadlines, as shown in Table 1.1:

	J'_1	J'_2	J'_3	J'_4
C'_i	2	4	3	5
D'_i	5	9	10	16

Table 1.1 Task set ordered by deadline.

Then applying equation (??) we have:

$$\begin{aligned}
 f'_1 &= C'_1 = 2 \\
 f'_2 &= f'_1 + C'_2 = 6 \\
 f'_3 &= f'_2 + C'_3 = 9 \\
 f'_4 &= f'_3 + C'_4 = 14
 \end{aligned}$$

Since each finishing time is less than the corresponding deadline, the task set is schedulable by EDD.

- 3.2 The algorithm for finding the maximum lateness of a task set scheduled by the EDD algorithm is shown in Figure 1.1.
- 3.3 The scheduling tree constructed by the Bratley's algorithm for the following set of non-preemptive tasks is illustrated in Figure 1.2.

	J_1	J_2	J_3	J_4
a_i	0	4	2	6
C_i	6	2	4	2
D_i	18	8	9	10

Table 1.2 Task set parameters for the Bratley's algorithm.

- 3.4 The schedule found by the Spring algorithm on the scheduling tree developed in the previous exercise with the heuristic function $H = a + C + D$ is $\{J_2, J_3, J_4, J_1\}$ which is unfeasible, since J_3 and J_4 miss their deadlines. Noticed that the same schedule is found with $H = D$, whereas the feasible solution is found with $H = a + D$.

```

Algorithm: EDD- $L_{max}(\mathcal{J})$ 
{
   $L_{max} = -D_n$ ;
   $f_0 = 0$ ;
  for (each  $J_i \in \mathcal{J}$ ) {
     $f_i = f_{i-1} + C_i$ ;
     $L_i = f_i - D_i$ ;
    if ( $L_i > L_{max}$ )  $L_{max} = L_i$ ;
  }
  return( $L_{max}$ );
}

```

Figure 1.1 Algorithm for finding the maximum lateness of a task set scheduled by EDD.




Figure 1.2 Scheduling tree constructed by the Bratley's algorithm for the task set shown in Table 1.2.

3.5 The precedence graph is shown in Figure 1.3.

Applying the transformation algorithm by Chetto and Chetto we get the parameters shown in Table 1.3.

So the schedule produced by EDF will be: $\{B, A, D, C, E, F, G\}$.




Figure 1.3 Precedence graph for Exercise 3.5.

	C_i	r_i	r^*_i	d_i	d^*_i
A	2	0	0	25	20
B	3	0	0	25	15
C	3	0	3	25	23
D	5	0	3	25	20
E	1	0	6	25	25
F	2	0	8	25	25
G	5	0	8	25	25

Table 1.3 Task set parameters modified by the Chetto and Chetto's algorithm.

SOLUTIONS FOR CHAPTER 4

4.1 The processor utilization factor of the task set is

$$U = \frac{2}{6} + \frac{2}{8} + \frac{2}{12} = 0.75$$

and considering that for three tasks the utilization least upper bound is

$$U_{lub}(3) = 3(2^{1/3} - 1) \simeq 0.78$$

from the Liu and Layland test, since $U \leq U_{lub}$, we can conclude that the task set is schedulable by RM, as shown in Figure 1.4.




Figure 1.4 Schedule produced by Rate Monotonic for the task set of Exercise 4.1.

4.2 The processor utilization factor of the task set is

$$U = \frac{3}{5} + \frac{1}{8} + \frac{1}{10} = 0.825$$

which is greater than $U_{lub}(3)$. Hence, we cannot verify the feasibility with the Liu and Layland test. Using the Hyperbolic Bound, we have that:

$$\prod_{i=1}^n (U_i + 1) = 1.98$$

which is less than 2. Hence, we can conclude that the task set is schedulable by RM, as shown in Figure 1.5.




Figure 1.5 Schedule produced by Rate Monotonic for the task set of Exercise 4.2.

4.3 Applying the Liu and Layland test we have that

$$U = \frac{1}{4} + \frac{2}{6} + \frac{3}{10} = 0.88 > 0.78$$

so we cannot say anything. With the Hyperbolic Bound we have that

$$\prod_{i=1}^n (U_i + 1) = 2.16 > 2$$

so we cannot say anything. Applying the Response Time Analysis we have to compute the response times and verify that they are less than or equal to the relative deadlines (which in this case are equal to periods). Hence, we have:

$$R_1 = C_1 = 1$$

So τ_1 does not miss its deadline. For τ_2 we have:

$$\begin{aligned} R_2^{(0)} &= \sum_{j=1}^2 C_j = C_1 + C_2 = 3 \\ R_2^{(1)} &= C_2 + \left\lceil \frac{R_2^{(0)}}{T_1} \right\rceil C_1 = 2 + \left\lceil \frac{3}{4} \right\rceil 1 = 3 \end{aligned}$$

So $R_2 = 3$, meaning that τ_2 does not miss its deadline. For τ_3 we have:

$$\begin{aligned} R_3^{(0)} &= \sum_{j=1}^3 C_j = C_1 + C_2 + C_3 = 6 \\ R_3^{(1)} &= C_3 + \left\lceil \frac{R_3^{(0)}}{T_1} \right\rceil C_1 + \left\lceil \frac{R_3^{(0)}}{T_2} \right\rceil C_2 = 3 + \left\lceil \frac{6}{4} \right\rceil 1 + \left\lceil \frac{6}{6} \right\rceil 2 = 7 \\ R_3^{(2)} &= 3 + \left\lceil \frac{7}{4} \right\rceil 1 + \left\lceil \frac{7}{6} \right\rceil 2 = 9 \\ R_3^{(3)} &= 3 + \left\lceil \frac{9}{4} \right\rceil 1 + \left\lceil \frac{9}{6} \right\rceil 2 = 10 \\ R_3^{(4)} &= 3 + \left\lceil \frac{10}{4} \right\rceil 1 + \left\lceil \frac{10}{6} \right\rceil 2 = 10 \end{aligned}$$

So $R_3 = 10$, meaning that τ_3 does not miss its deadline. Hence we can conclude that the task set is schedulable by RM, as shown in Figure 1.6.




Figure 1.6 Schedule produced by Rate Monotonic for the task set of Exercise 4.3.

4.4 Applying the Response Time Analysis, we can easily verify that $R_3 = 10$ (see the solution of the previous exercise), hence the task set is not schedulable by RM.

4.5 Since

$$U = \frac{1}{4} + \frac{2}{6} + \frac{3}{8} = 0.96 < 1$$

the task set is schedulable by EDF, as shown in Figure 1.7.




Figure 1.7 Schedule produced by EDF for the task set of Exercise 4.5.

4.6 Applying the processor demand criterion, we have to verify that

$$\forall L \in \mathcal{D} \quad \sum_{i=1}^n \left\lfloor \frac{L + T_i - D_i}{T_i} \right\rfloor C_i \leq L.$$

where

$$\mathcal{D} = \{d_k \mid d_k \leq \min(L^*, H)\}.$$

For the specific example, we have

$$\begin{aligned} U &= \frac{2}{6} + \frac{2}{8} + \frac{4}{12} = \frac{11}{12} \\ L^* &= \frac{\sum_{i=1}^n (T_i - D_i) U_i}{1 - U} = 32 \\ H &= \text{lcm}(6, 8, 12) = 24. \end{aligned}$$

Hence, the set of checking points is given by $\mathcal{D} = \{4, 5, 8, 11, 12, 17, 20, 23\}$. Since the demand in these intervals is $\{2, 4, 8, 10, 12, 14, 20, 22\}$ we can conclude that the task set is schedulable by EDF. The resulting schedule is shown in Figure 1.8.

4.7 Applying the Response Time Analysis, we have to start by computing the response time of task τ_2 , which is the one with the shortest relative deadline, and hence the highest priority:

$$R_2 = C_2 = 2.$$




Figure 1.8 Schedule produced by EDF for the task set of Exercise 4.6.

So τ_2 does not miss its deadline. For τ_1 we have:

$$R_1^{(0)} = \sum_{j=1}^2 C_j = C_1 + C_2 = 4$$

$$R_1^{(1)} = C_1 + \left\lceil \frac{R_1^{(0)}}{T_2} \right\rceil C_2 = 2 + \left\lceil \frac{4}{8} \right\rceil 2 = 4$$

So $R_1 = 4$, meaning that τ_1 does not miss its deadline. For τ_3 we have:

$$R_3^{(0)} = \sum_{j=1}^3 C_j = C_1 + C_2 + C_3 = 8$$

$$R_3^{(1)} = C_3 + \left\lceil \frac{R_3^{(0)}}{T_2} \right\rceil C_2 + \left\lceil \frac{R_3^{(0)}}{T_1} \right\rceil C_1 = 4 + \left\lceil \frac{8}{8} \right\rceil 2 + \left\lceil \frac{8}{6} \right\rceil 2 = 10$$

And since $R_3^{(1)} > D_3$, we can conclude that the task set is not schedulable by DM. The resulting schedule is shown in Figure 1.9.




Figure 1.9 Schedule produced by Deadline Monotonic for the task set of Exercise 4.7.

SOLUTIONS FOR CHAPTER 5

- 5.1 Since the Sporadic Server behaves in the worst case as a periodic task, it can be guaranteed with the condition derived for the Priority Exchange Server. Hence, considering the result expressed by equation (??) a periodic task set can be guaranteed together with a Sporadic Server under RM, if

$$U_p \leq n \left[\left(\frac{2}{U_s + 1} \right)^{1/n} - 1 \right].$$

Inverting this relation we have:

$$U_s \leq 2 \left(\frac{U_p}{n} + 1 \right)^{-n} - 1$$

and considering that $n = 2$ and $U_p = 0.45$, we have that

$$U_{s_{max}} = 0.33$$

- 5.2 Considering the result expressed by equation (??), a periodic task set can be guaranteed together with a Deferrable Server under RM, if

$$U_p \leq n \left[\left(\frac{U_s + 2}{2U_s + 1} \right)^{1/n} - 1 \right].$$

Inverting this relation we have:

$$U_s \leq \frac{2 - K}{2K - 1},$$

where $K = (U_p/n + 1)^n$. And considering that $n = 2$ and $U_p = 0.45$, we have $K = 1.5$, and $U_{s_{max}} = 0.25$.

- 5.3 From exercise 5.??, we know that the maximum server utilization that can be assigned to a Polling Server to guarantee the periodic task set is $U_{s_{max}} = 0.33$. So, by setting $T_s = 6$ (intermediate priority) and $C_s = 2$, we satisfy the constraints. The resulting schedule is illustrated in Figure 1.10.
- 5.4 A Sporadic Server can be guaranteed with the same method used for the Polling Server. So, using the same parameters computed before ($C_s = 2$ and $T_s = 6$) we have the schedule shown in Figure 1.11.




Figure 1.10 Schedule produced by Rate Monotonic and Polling Server for the task set of Exercise 5.3.




Figure 1.11 Schedule produced by Rate Monotonic and Sporadic Server for the task set of Exercise 5.4.

- 5.5 From exercise 5.??, we know that the maximum server utilization that can be assigned to a Deferrable Server to guarantee the periodic task set is $U_{s_{m.a.x}} = 0.25$. So, by setting $T_s = 4$ (maximum priority) and $C_s = 1$, we satisfy the constraints. The resulting schedule is illustrated in Figure 1.12.
- 5.6 The resulting schedule is illustrated in Figure 1.13.

SOLUTIONS FOR CHAPTER 6




Figure 1.12 Schedule produced by Rate Monotonic and Deferrable Server for the task set of Exercise 5.5.




Figure 1.13 Schedule produced by Rate Monotonic and Sporadic Server for the task set of Exercise 5.6.

- 6.1 For any dynamic server we must have $U_p + U_s \leq 1$, hence, considering that $U_p = 2/3$, the maximum server utilization that can be assigned to a Dynamic Sporadic Server is:

$$U_s = 1 - U_p = 1/3.$$

- 6.2 The deadlines computed by the server for the aperiodic jobs result: $d_1 = a_1 + T_s = 7$, $d_2 = d_1 + T_s = 13$, and $d_3 = a_3 + T_s = 21$. The resulting schedule produced by EDF + DSS is illustrated in Figure 1.14.

- 6.3 The deadlines computed by the server for the aperiodic jobs are: $d_1 = a_1 + C_1/U_s = 10$, $d_2 = d_1 + C_2/U_s = 13$, and $d_3 = a_3 + C_3/U_s = 18$. The resulting schedule produced by EDF + TBS is illustrated in Figure 1.15.




Figure 1.14 Schedule produced by EDF + DDS for the task set of Exercise 6.2.




Figure 1.15 Schedule produced by EDF + TBS for the task set of Exercise 6.3.

6.4 The events handled by the CBS are:

time	event	action
$t = 1$	arrival	$c_s = Q_s, d_s = a_1 + T_s = 7$
$t = 4$	$c_s = 0$	$c_s = Q_s, d_s = d_s + T_s = 13$
$t = 5$	arrival	enqueue request
$t = 11$	$c_s = 0$	$c_s = Q_s, d_s = d_s + T_s = 19$
$t = 15$	arrival	$c_s = Q_s, d_s = a_3 + T_s = 21$

The resulting schedule produced by EDF + CBS is illustrated in Figure 1.16.

6.5 The deadlines computed by the server are:

$$\begin{aligned}
 d_1^{(0)} &= a_1 + C_1/U_s = 10 \\
 d_1^{(1)} &= f_1^{(0)} = 8
 \end{aligned}$$




Figure 1.16 Schedule produced by EDF + CBS for the task set of Exercise 6.4.

$$d_2^{(0)} = d_2^{(0)} + C_2/U_s = 13$$

$$d_2^{(1)} = f_2^{(0)} = 11$$

$$d_3^{(0)} = a_3 + C_3/U_s = 18$$

$$d_3^{(1)} = f_3^{(0)} = 16$$

The resulting schedule produced by EDF + TB(1) is illustrated in Figure 1.17.




Figure 1.17 Schedule produced by EDF + TB(1) for the task set of Exercise 6.5.

6.6 The deadlines computed by the server are:

$$d_1^{(0)} = a_1 + C_1/U_s = 10$$

$$d_1^{(1)} = f_1^{(0)} = 8$$

$$d_1^{(2)} = f_1^{(1)} = 4$$

$$d_1^{(3)} = f_1^{(2)} = 3$$

$$d_2^{(0)} = d_2^{(0)} + C_2/U_s = 13$$

$$d_2^{(1)} = f_2^{(0)} = 11$$

$$d_2^{(2)} = f_2^{(1)} = 9$$

$$d_2^{(3)} = f_2^{(2)} = 6$$

$$d_3^{(0)} = a_3 + C_3/U_s = 18$$

$$d_3^{(1)} = f_3^{(0)} = 16$$

The resulting schedule produced by EDF + TB* is illustrated in Figure 1.18.




Figure 1.18 Schedule produced by EDF + TB* for the task set of Exercise 6.6.

6.7 The resulting schedule is illustrated in Figure 1.19.

SOLUTIONS FOR CHAPTER 7

7.1 Applying equation (??), we have verify that:

$$\forall i, 1 \leq i \leq n, \quad \sum_{k=1}^i \frac{C_k}{T_k} + \frac{B_i}{T_i} \leq i(2^{1/i} - 1).$$

So we have:

$$\frac{C_1 + B_1}{T_1} = \frac{9}{10} < 1$$




Figure 1.19 Schedule produced by EDF+T B₁+T B₂ for the task set of Exercise 6.7.

$$\frac{C_1}{T_1} + \frac{C_2 + B_2}{T_2} = \frac{4}{10} + \frac{6}{15} = 0.8 < 0.83$$

$$\frac{C_1}{T_1} + \frac{C_2}{T_2} + \frac{C_3}{T_3} = \frac{4}{10} + \frac{3}{15} + \frac{4}{20} = 0.8 > 0.78$$

So we cannot say anything about feasibility. By applying the response time analysis we have to verify that:

$$\forall i, 1 \leq i \leq n, \quad R_i \leq D_i$$

where

$$R_i = C_i + B_i + \sum_{k=1}^{i-1} \left\lceil \frac{R_i}{T_k} \right\rceil C_k$$

So we have:

$$R_1 = C_1 + B_1 = 9 < 10$$

$$R_2^{(0)} = C_1 + C_2 + B_2 = 10$$

$$R_2^{(1)} = C_2 + B_2 + \left\lceil \frac{10}{10} \right\rceil 4 = 10 < 15$$

$$R_3^{(0)} = C_1 + C_2 + C_3 = 11$$

$$R_3^{(1)} = C_3 + \left\lceil \frac{11}{10} \right\rceil 4 + \left\lceil \frac{11}{15} \right\rceil 3 = 15$$

$$R_3^{(2)} = C_3 + \left\lceil \frac{15}{10} \right\rceil 4 + \left\lceil \frac{15}{15} \right\rceil 3 = 15 < 20$$

Hence we can conclude that the task set is schedulable by RM.

- 7.2 Using the Priority Inheritance Protocol, a task τ_i can be blocked at most for one critical section by each lower priority task. Moreover, a critical section can block τ_i only if it belongs to a task with lower priority and it is shared with τ_i (direct blocking) or with higher priority tasks (push-through blocking). Finally, we have to consider that two critical sections cannot block a task if they are protected by the same semaphore or they belong to the same task. Hence, if X_i denotes the critical section X belonging to task τ_i we have that:
- Task τ_1 can only experience direct blocking (since there are no tasks with higher priority) and the set of critical sections that can potentially block it is $\{A_2, A_3, C_3\}$. It can be blocked at most for the duration of two critical sections in this set. Thus, the maximum blocking time is given by the sum of the two longest critical sections in this set. In this case, however, note that the longest critical sections are A_3 and C_3 , which belong to the same task, hence they cannot be selected together. Hence, the maximum blocking time for τ_1 is $B_1 = d_{(A_2)} + d_{(C_3)} = 7$.
 - Task τ_2 can experience direct blocking on A_3 and B_3 and push-through blocking on A_3 and C_3 . Hence, the set of critical sections that can potentially block τ_2 is $\{A_3, B_3, C_3\}$. It can be blocked at most for the duration of one critical section in this set. Thus, the maximum blocking time is given by the longest critical section in this set, that is C_3 . Hence, we have $B_2 = d_{(C_3)} = 5$.
 - Task τ_3 cannot be blocked, because it is the task with the lowest priority (it can only be preempted by higher priority tasks). Hence, we have $B_3 = 0$.
- 7.3 Using the Priority Ceiling Protocol, a task τ_i can be blocked at most for one critical section during its execution. The set of critical sections that can potentially block τ_i is the same as that computed for the Priority Inheritance Protocol. Hence, if X_i denotes the critical section X belonging to task τ_i we have that:
- The set of critical sections that can potentially block τ_1 is $\{A_2, A_3, C_3\}$. Hence, the maximum blocking time for τ_1 is $B_1 = d_{(C_3)} = 5$.
 - The set of critical sections that can potentially block τ_2 is $\{A_3, B_3, C_3\}$. Hence, the maximum blocking time for τ_2 is $B_2 = d_{(C_3)} = 5$.
 - Task τ_3 cannot be blocked, because it is the task with the lowest priority (it can only be preempted by higher priority tasks). Hence, we have $B_3 = 0$.
- 7.4 The maximum blocking time for τ_2 is given by a push-through blocking on C_3 . This means that, for this to happen, τ_3 must start first and must enter its critical section C_3 . Then, τ_1 must preempt τ_3 , so that τ_3 can inherit the highest priority to prevent τ_2 to execute. The situation is illustrated in Figure 1.20.




Figure 1.20 Schedule produced by RM + PIP for the task set of Exercise 7.4.

7.5 To compute the maximum blocking time under the Priority Inheritance Protocol we reason as follows.

- The set of critical sections that can potentially block τ_1 is $\{C_2, B_3, D_3, E_3, A_4, C_4, E_4\}$. Among these, we have to select the three longest ones, one for each lower priority task. Note that, if we select C_2 and E_3 , we cannot select E_4 (which is the longest of τ_4) because E has been already selected for τ_3 , and we cannot select C_4 for the same reason. So, we have to select A_4 . Hence, the maximum blocking time for τ_1 is $B_1 = d_{(C_2)} + d_{(E_3)} + d_{(A_4)} = 26$.
- Task τ_2 can experience direct blocking on C_4 and push-through blocking on B_3, E_3, A_4, C_4 , and E_4 . Hence, the set of critical sections that can potentially block τ_2 is $\{B_3, E_3, A_4, C_4, E_4\}$. It can be blocked at most for the duration of two critical sections in this set. Thus, we have $B_2 = d_{(E_3)} + d_{(C_4)} = 21$. Note that E_3 and E_4 cannot block τ_2 together.
- Task τ_3 can experience direct blocking on E_4 and push-through blocking on A_4, C_4 , and E_4 . Hence, the set of critical sections that can potentially block τ_3 is $\{A_4, C_4, E_4\}$. It can be blocked at most for the duration of one critical section in this set. Thus, we have $B_3 = d_{(E_4)} = 10$.
- Task τ_4 cannot be blocked, because it is the task with the lowest priority (it can only be preempted by higher priority tasks). Hence, we have $B_4 = 0$.

7.6 The sets of critical sections that can cause blocking under the Priority Ceiling Protocol are the same as those derived in the previous exercise for the Priority Inheritance Protocol. The only difference is that under the Priority Ceiling Protocol each task can only be blocked for the duration of a single critical section. Hence, we have:

- The set of critical sections that can potentially block τ_1 is $\{C_2, B_3, D_3, E_3, A_4, C_4, E_4\}$. Hence, the maximum blocking time for τ_1 is $B_1 = d_{(E_3)} = 13$.

- The set of critical sections that can potentially block τ_2 is $\{B_3, E_3, A_4, C_4, E_4\}$. Hence, the maximum blocking time for τ_2 is $B_2 = d_{(E_3)} = 13$.
- The set of critical sections that can potentially block τ_3 is $\{A_4, C_4, E_4\}$. Hence, the maximum blocking time for τ_3 is $B_3 = d_{(E_4)} = 10$.
- Task τ_4 cannot be blocked, because it is the task with the lowest priority (it can only be preempted by higher priority tasks). Hence, we have $B_4 = 0$.

7.7 The maximum blocking time for τ_2 is given by a push-through blocking on C_4 and E_3 . This means that, for this to happen, τ_4 must start first and must enter its critical section C_4 . Then, τ_3 must preempt τ_4 , entering E_3 . Now, when τ_1 arrives, it experiences a chained blocking when entering C_1 and E_1 , which are both locked. The situation is illustrated in Figure 1.21.




Figure 1.21 Schedule produced by RM + PIP for the task set of Exercise 7.7.

7.8 If tasks are assigned decreasing preemption levels as $\pi_1 = 3$, $\pi_2 = 2$, and $\pi_3 = 1$, the resource ceilings have the values shown in Table 1.4.

	$C_R(3)$	$C_R(2)$	$C_R(1)$	$C_R(0)$
A	0	1	2	3
B	0	0	0	2
C	-	0	2	3

Table 1.4 SRP resource ceilings for Exercise 7.8.